

Person wants to ZOOM  
in DEEP into the  
Mandelbrot Set...



...but they are limited by  
their weak device...



Distributed computing  
will save the day.



User Assembles a Request


```
{  
  START: [0.29340, -1.3941],  
  ZOOM: "8x",  
  NUM_FRAMES: 128,  
  FPS: 20,  
  THEME: "Galaxy",  
  FORMAT: "mp4"  
}
```

### CLUSTER LEADER (PHASE I)


Sequence generation "task" is broken  
down into the following steps:

1. Compute the sets of (real, imaginary)  
coordinate pairs that define all of the  
frames needed to construct a request.
2. Check if frames have been generated in  
the past before starting computation.
3. Assemble an array of promised 'futures'  
that represent atomic units of workload.  
(@ray.remote)
4. Execute the futures using Ray.io to  
gather all available GPU/CPU cores and  
summon them (@ray.get)


WORKER NODES



Each worker executes  
pyopencil instructions  
that accommodate CPU,  
GPU, and many other  
optimizations



INFINITE inductive bias makes  
this an INFINITELY scalable  
problem.




Work is solely a function of the number of  
pixels for which the iterative convergence/  
divergence process is preformed. Frames,  
grids, sequences, are all just recursively  
built up from pixels.

...


Each workers fulfills  
promises by processing N  
frames.

Frames encoded as b64 strings.


START: 0.29340Re, -1.3941 Im



ZOOM: 8x




ZOOM: 16x



### CLUSTER LEADER (PHASE II)

1. Persist any novel frames.
2. Stitch everything together and serve  
the users request!



Monitor Nodes using dialogue. Launch Dashboard  
for performance stats.

